

🔗 SQL OPERATORS — REAL TIME USAGE (Practical Point of View)

◆ □ Arithmetic Operators

*(+, -, , /, %)

👉 Used very rarely in real-time queries directly — mostly in reports, calculations, and dashboards.

◆ Real-time Use:

| Domain | Example | Query |
|---------|------------------------------------|--------------------------------------------------------------|
| College | Add 5 marks bonus for each student | SELECT student_name, marks + 5 AS total FROM student; |
| Payroll | Calculate annual salary | SELECT emp_name, salary * 12 AS yearly_salary FROM employee; |
| Finance | Calculate profit | SELECT revenue - cost AS profit FROM accounts; |

🌐 Usage Frequency: ◆ Medium (used in calculations, reporting)

◆ □ Comparison Operators

(=, !=, <>, >, <, >=, <=)

👉 These are the most common — used in WHERE conditions, JOIN conditions, filters etc.

◆ Real-time Use:

| Domain | Example | Query |
|------------|-----------------------------------|-------------------------------|
| College | Get students with marks ≥ 60 | WHERE marks ≥ 60 |
| HR | Employees with salary < 50000 | WHERE salary < 50000 |
| Banking | Customers with balance $\neq 0$ | WHERE balance $\neq 0$ |
| E-Commerce | Orders delivered = 'YES' | WHERE delivery_status = 'YES' |

🔗 Usage Frequency: ⬤ Very High (in 90% of queries)

◆ Logical Operators

(AND, OR, NOT)

👉 Used to combine multiple conditions — practically used in every project.

◆ Real-time Use:

| Domain | Example | Query |
|---------|-------------------------------------|----------------------------------------------|
| College | Students in CSE with marks > 80 | WHERE dept_name = 'CSE' AND marks > 80 |
| Banking | Account active OR balance > 10000 | WHERE status = 'ACTIVE' OR balance > 10000 |
| HR | NOT from HR Department | WHERE NOT dept = 'HR' |

🔍 Usage Frequency: ⬤ Very High (used in almost every WHERE clause)

💎 4️⃣ Special Operators — Most Important Real-Time Usage

ఈ సెక్షన్లోకుభగా వాడతారు real-time లో 📌

✅ (A) BETWEEN

👉 Used for range checking (numeric or date ranges).

| Domain | Example | Query |
|---------|---------------------------------------|------------------------------------------------------|
| College | Students with marks 70 to 90 | WHERE marks BETWEEN 70 AND 90 |
| Finance | Transactions between 1-Jan and 31-Jan | WHERE txn_date BETWEEN '2025-01-01' AND '2025-01-31' |
| HR | Employees with salary range 40k–80k | WHERE salary BETWEEN 40000 AND 80000 |

🔍 Usage Frequency: ⬤ High

💡 Used mainly in date filters and numeric range reports.

✅ (B) IN / NOT IN

👉 Used to check multiple values without using multiple ORs.

| Domain | Example | Query |
|---------|-------------------------------|---------------------------------------|
| College | Students in CSE or ECE | WHERE dept_name IN ('CSE', 'ECE') |
| HR | Exclude interns | WHERE role NOT IN ('Intern') |
| Banking | Transactions of certain types | WHERE txn_type IN ('CREDIT', 'DEBIT') |

🔍 Usage Frequency: ⬤ Very High

💡 Used heavily in filtering, subqueries, and reports.

✓ (C) LIKE / ILIKE (PostgreSQL specific)

👉 Used for pattern-based search (string matching).

| Domain | Example | Query |
|---------|-------------------------------------|-------------------------------|
| College | Students whose name starts with 'R' | WHERE student_name LIKE 'R%' |
| HR | Employees whose name ends with 'a' | WHERE emp_name LIKE '%a' |
| Banking | Account numbers containing '123' | WHERE account_no LIKE '%123%' |

🔄 Usage Frequency: ● Very High

💡 PostgreSQL లో ILIKE నూడా ఉంది → case-insensitive search కోసం.

👉 ILIKE అంటే 'r%' or 'R%' రెండూ match అవుతాయి.

✓ (D) IS NULL / IS NOT NULL

👉 To find missing or undefined data (very common in real-time).

| Domain | Example | Query |
|---------|------------------------------------|-----------------------|
| College | Students with fees not entered | WHERE fees IS NULL |
| HR | Employees whose salary not updated | WHERE salary IS NULL |
| Finance | Transactions where remarks missing | WHERE remarks IS NULL |



🔄 Usage Frequency: ● Very High


💡 Used a lot in data quality checks and NULL validations.


✓ (E) EXISTS / NOT EXISTS


👉 Used with subqueries to check record existence.

| Domain | Example | Query |
|---------|---------------------------------------|----------------------------------------------------------------------------------------------------------------------|
| College | Departments with at least one student | sql SELECT dept_name FROM department d WHERE EXISTS (SELECT 1 FROM student s WHERE s.dept_id = d.dept_id); |
| HR | Employees having assigned projects | sql WHERE EXISTS (SELECT 1 FROM project WHERE emp_id = e.emp_id); |

 **Usage Frequency:**  *Medium-High*

 Used in **joins optimization** and **business validations**.

 ** DISTINCT / DISTINCT ON**

 Used to remove duplicates in reports or summary queries.

| Domain | Example | Query |
|------------|---------------------------------------------|-------------------------------------------------------------------------------|
| College | Unique department names | SELECT DISTINCT dept_name FROM student; |
| HR | Unique job roles | SELECT DISTINCT job_role FROM employee; |
| PostgreSQL | Get one record per department (DISTINCT ON) | SELECT DISTINCT ON (dept_name) * FROM student ORDER BY dept_name, marks DESC; |

 **Usage Frequency:**  *Medium (in reports & analytics)*

◆ **Operator Usage Summary (Real Time Importance)**

| Operator Type | Examples | Real-Time Usage | Importance |
|------------------------|---------------------|-----------------|------------|
| Arithmetic | +, -, *, /, % | Moderate | ★ ★ |
| Comparison | =, <, >, <=, >=, <> | Very High | ★ ★ ★ ★ ★ |
| Logical | AND, OR, NOT | Very High | ★ ★ ★ ★ ★ |
| BETWEEN | Range checking | High | ★ ★ ★ ★ |
| IN / NOT IN | List matching | Very High | ★ ★ ★ ★ ★ |
| LIKE / ILIKE | Pattern search | Very High | ★ ★ ★ ★ ★ |
| IS NULL / IS NOT NULL | Missing data check | Very High | ★ ★ ★ ★ ★ |
| EXISTS / NOT EXISTS | Subquery validation | High | ★ ★ ★ ★ |
| DISTINCT / DISTINCT ON | Unique data | Medium | ★ ★ ★ |

🧠 **Real-Time Project Usage Summary**

| Project Type | Most Common Operators | Usage Example |
|-----------------|----------------------------|-------------------------------|
| College ERP | BETWEEN, IN, LIKE, IS NULL | Student reports, fee analysis |
| Banking | BETWEEN, IN, EXISTS | Transactions by date range |
| HRMS | IN, BETWEEN, IS NOT NULL | Employee filters |
| E-Commerce | LIKE, BETWEEN, IN | Product search, sales filter |
| Finance / Audit | BETWEEN, EXISTS, IS NULL | Monthly statements |

✓ **Conclusion**

- > In ****real-time SQL****, you will use mainly:
- > - ****Comparison operators**** (=, >, <, <=, >=)
- > - ****Logical operators**** (AND, OR, NOT)
- > - ****Special operators****:
 - > ♦ `BETWEEN`
 - > ♦ `IN` / `NOT IN`
 - > ♦ `LIKE` / `ILIKE`
 - > ♦ `IS NULL` / `IS NOT NULL`
 - > ♦ `EXISTS`

☞ ఇవి 90% SQL queriesలో ఉంటాయి —

Joins, subqueries, group by, having అన్నిటిలో కూడా వీటినే ఉపయోగిస్తారు.

✿ SQL OPERATORS — BASIC INTERVIEW QUESTIONS & ANSWERS

| No | Question | Answer |
|----|-------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 | What are SQL Operators? | Operators are special symbols or keywords used to perform operations on data values — like comparison, arithmetic, logical, and pattern matching. |
| 2 | How many types of operators are there in SQL? | Mainly 5 types: 1 Arithmetic 2 Comparison 3 Logical 4 Special (LIKE, BETWEEN, IN, IS NULL) 5 Set operators (UNION, INTERSECT, EXCEPT). |
| 3 | What is the difference between = and == in SQL? | SQL uses only one equal sign = for comparison. == is not valid in SQL. |
| 4 | What is the use of the BETWEEN operator? | To check whether a value lies between two given values (range checking). Example: WHERE marks BETWEEN 60 AND 90. |
| 5 | What is the difference between BETWEEN and IN? | BETWEEN → for continuous ranges (e.g., 60–90). IN → for specific values (e.g., 'CSE', 'ECE'). |
| 6 | What is the use of the LIKE operator? | It is used for pattern matching in text columns. Example: WHERE name LIKE 'R%'. |
| 7 | What is the difference between LIKE and ILIKE? | LIKE → case-sensitive (SQL Server, MySQL). ILIKE → case-insensitive (PostgreSQL only). |
| 8 | What is the use of IS NULL and IS NOT NULL? | Used to check whether a column contains NULL (missing) values. Example: WHERE salary IS NULL. |
| 9 | What is the use of the IN operator? | To check if a value matches any value in a list. Example: WHERE dept IN ('CSE', 'ECE', 'EEE'). |
| 10 | What is the difference between IN and EXISTS? | IN checks a list of values; EXISTS checks whether a subquery returns any rows. EXISTS is faster for large datasets. |

| No | Question | Answer |
|----|----------------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| 11 | What does the DISTINCT operator do? | Removes duplicate rows from the result set. Example: <code>SELECT DISTINCT dept_name FROM student;</code> |
| 12 | Is DISTINCT supported in PostgreSQL? | Yes, and it also supports DISTINCT ON (column) for advanced distinct filtering. |
| 13 | What is the use of arithmetic operators in SQL? | To perform mathematical calculations. Example: <code>salary * 12</code> gives yearly salary. |
| 14 | What is the difference between AND and OR? | AND → All conditions must be true. OR → At least one condition must be true. |
| 15 | What is the purpose of the NOT operator? | It reverses the condition result. Example: <code>WHERE NOT dept='CSE'</code> gives all except CSE. |
| 16 | What is the difference between = and <>? | = is used for equality, <> (or !=) for inequality. Example: <code>WHERE dept <> 'CSE'</code> . |
| 17 | Can we use BETWEEN with dates? | Yes. Example: <code>WHERE join_date BETWEEN '2024-01-01' AND '2024-12-31'</code> . |
| 18 | Can we use LIKE with numbers? | Technically yes (after converting numbers to text), but normally used only with character/text columns. |
| 19 | Which operator is used for multiple OR conditions? | IN operator is used instead of multiple ORs. Example: <code>WHERE dept IN ('CSE','ECE','MECH')</code> . |
| 20 | What happens if we use NOT IN with NULL values? | It may return no rows because NULL causes unknown comparison. Always use NOT EXISTS for safety. |

💡 Bonus Tips (asked in practical interviews):

| Topic | Interview Tip |
|-------|---------------|
|-------|---------------|

| Topic | Interview Tip |
|----------|--------------------------------------------------------------------------|
| BETWEEN | Always inclusive — means both lower and upper limits are included. |
| LIKE | % → any number of characters, _ → single character. |
| DISTINCT | Slows down queries if used unnecessarily; avoid unless duplicates exist. |
| EXISTS | Used for performance optimization in subqueries. |
| NULL | Comparison with = or <> does not work; always use IS NULL. |

=====

=====

=====

🔗 SQL OPERATORS — INTERMEDIATE & REAL-TIME QUESTIONS WITH ANSWERS

| No. | Question | Answer / Explanation |
|-----|---------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 | What is the difference between IN and EXISTS operator in SQL? | ◆ IN checks whether a value is in a list of constants or subquery results. ◆ EXISTS checks if a subquery returns at least one record. ✓ EXISTS is faster for large datasets because it stops searching after finding the first match. |
| 2 | What is the difference between BETWEEN and < / > operators? | BETWEEN checks a range (inclusive). Example: BETWEEN 10 AND 20 = >=10 AND <=20. < / > are used for individual comparisons. |
| 3 | Can we use LIKE with numeric fields? | Usually not recommended, but possible by converting number to text. Example: WHERE CAST(emp_id AS TEXT) LIKE '12%'. |
| 4 | What is the use of ANY and ALL operators? | ◆ ANY → true if any value in subquery satisfies the condition. ◆ ALL → true only if all values in subquery satisfy the condition. Example: salary > ALL (SELECT salary FROM interns); |

| No. | Question | Answer / Explanation |
|-----|-------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------|
| 5 | Explain difference between = ANY and IN. | Both are equivalent — = ANY(subquery) works the same as IN(subquery). |
| 6 | Can LIKE operator use indexes in PostgreSQL? | Normally, LIKE '%abc' won't use index. But LIKE 'abc%' can use index if the column has a btree or text_pattern_ops index. |
| 7 | What is the difference between IS NULL and = NULL? | = NULL never returns true. You must use IS NULL or IS NOT NULL. |
| 8 | What is the difference between NOT IN and NOT EXISTS? | NOT IN fails if subquery returns any NULL . NOT EXISTS is safer and faster for large datasets. |
| 9 | Can you use DISTINCT with aggregate functions? | Yes <input checked="" type="checkbox"/> Example: COUNT(DISTINCT dept_id) — counts only unique departments. |
| 10 | What is the difference between DISTINCT and GROUP BY? | Both remove duplicates. DISTINCT is simple & direct. GROUP BY allows aggregation (SUM, COUNT, etc.). |
| 11 | What is operator precedence in SQL? | Order of execution: 1 Arithmetic (*, /, +, -) 2 Comparison (=, <, >) 3 Logical (NOT, AND, OR) |
| 12 | Can we combine AND, OR, and NOT together? | Yes <input checked="" type="checkbox"/> Example: WHERE (dept='CSE' OR dept='ECE') AND marks > 70 AND NOT name LIKE 'R%' |
| 13 | How does BETWEEN behave with strings? | Works alphabetically . Example: WHERE name BETWEEN 'A' AND 'M' → returns names from A to M. |
| 14 | What happens if NULL is compared using > or <? | It returns UNKNOWN . NULL can't be compared with operators. Use IS NULL instead. |

| No. | Question | Answer / Explanation |
|-----|-----------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 15 | Which is faster: IN or EXISTS? | For small lists → IN is fine. For subqueries → EXISTS is faster on large datasets. |
| 16 | What is the use of COALESCE() in conditions? | Replaces NULL values with default. Example: WHERE COALESCE(salary,0) > 5000. |
| 17 | What's the difference between = and IS? | = compares values. IS is used for conditions like IS NULL, IS TRUE, IS FALSE. |
| 18 | What is the difference between LIKE and SIMILAR TO? | LIKE uses simple wildcards (% , _). SIMILAR TO supports regex-like patterns. Example: `name SIMILAR TO '(A |
| 19 | What is the purpose of the CASE operator in SQL? | Conditional operator like IF-ELSE. Example: CASE WHEN marks >= 60 THEN 'Pass' ELSE 'Fail' END |
| 20 | Real-time: When would you use DISTINCT vs GROUP BY? | In dashboards/reports: <input checked="" type="checkbox"/> Use DISTINCT for unique departments. <input checked="" type="checkbox"/> Use GROUP BY to get aggregates (e.g. total salary per department). |

🔗 REAL-TIME PROJECT SCENARIOS

| Scenario | Operator Used | Example |
|--------------------------------------------|---------------|------------------------------------|
| Show all unique departments | DISTINCT | SELECT DISTINCT dept FROM student; |
| Find students with marks between 80 and 90 | BETWEEN | WHERE marks BETWEEN 80 AND 90; |
| Search all students whose names start with | LIKE | WHERE name LIKE 'S%'; |

| Scenario | Operator Used | Example |
|---------------------------------------------------------|---------------|------------------------------------------------------------------------|
| 'S' | | |
| Filter employees who are not assigned to any project | IS NULL | WHERE project_id IS NULL; |
| Show employees whose department is either HR or IT | IN | WHERE dept IN ('HR','IT'); |
| Show products sold in multiple regions | IN + DISTINCT | SELECT DISTINCT product_id FROM sales WHERE region IN ('East','West'); |
| Find employees whose salary is greater than any intern | ANY | WHERE salary > ANY (SELECT salary FROM interns); |
| Find employees whose salary is greater than all interns | ALL | WHERE salary > ALL (SELECT salary FROM interns); |
| Show all customers except those from inactive list | NOT IN | WHERE cust_id NOT IN (SELECT id FROM inactive_customers); |
| Display top 1 student from each department (PostgreSQL) | DISTINCT ON | SELECT DISTINCT ON (dept) * FROM student ORDER BY dept, marks DESC; |

🔗 Intermediate → Expert Quick Recap Table

| Operator | Description | Real-time Example |
|----------|-------------|-------------------|
|----------|-------------|-------------------|

| Operator | Description | Real-time Example |
|------------------------------|------------------------------------------------------------|------------------------------------------------------------------------------------|
| IN / NOT IN | Check if a value exists in a list | Filter selected departments |
| EXISTS / NOT EXISTS | Check for existence using a subquery | Validate related records (e.g. orders that have matching customers) |
| BETWEEN | Range filtering (inclusive) | Filter marks between 60 and 80, or dates within a month |
| LIKE / ILIKE | Pattern matching (ILIKE is case-insensitive in PostgreSQL) | Search names starting with 'A', emails ending in @gmail.com |
| IS NULL / IS NOT NULL | Check for missing or present values | Find unassigned records like employees without project ID |
| DISTINCT | Eliminate duplicate values | Generate list of unique cities or departments |
| ANY / ALL | Compare with one or all values from a subquery | Check if salary > ANY intern salary or > ALL intern salaries |
| CASE | Conditional logic (like IF-ELSE) | Assign grades (CASE WHEN marks >= 60 THEN 'Pass' ELSE 'Fail') or calculate bonuses |

=====

=====


| No | Question | Explanation / Optimization Tip |
|----|-----------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 | What is the main difference between EXISTS and JOIN, and when to use each? | <ul style="list-style-type: none"> ◆ JOIN retrieves and merges rows. ◆ EXISTS checks existence (True/False). <input checked="" type="checkbox"/> Use EXISTS for filtering; JOIN when you need actual data. ✂ Example: SELECT e.emp_name FROM employee e WHERE EXISTS (SELECT 1 FROM department d WHERE e.dept_id = d.dept_id); |
| 2 | Difference between IN and EXISTS in terms of performance? | <ul style="list-style-type: none"> ◆ IN loads all values first — better for small sets. ◆ EXISTS exits early — better for large data. <input checked="" type="checkbox"/> Replace IN with EXISTS for subqueries returning many rows. |
| 3 | How can ANY and ALL operators be optimized? | <ul style="list-style-type: none"> ◆ > ANY (subquery) = greater than min. ◆ > ALL (subquery) = greater than max. <input checked="" type="checkbox"/> Use aggregate functions directly. ✂ Example: SELECT emp_name FROM employee WHERE salary > (SELECT MAX(salary) FROM interns); |
| 4 | Use of CASE in performance tuning? | <ul style="list-style-type: none"> ◆ Replaces multiple OR/IF conditions. <input checked="" type="checkbox"/> Ideal for dashboards or derived status fields. ✂ Example: SELECT emp_id, CASE WHEN salary > 80000 THEN 'HIGH' WHEN salary > 50000 THEN 'MEDIUM' ELSE 'LOW' END AS salary_level FROM employee; |
| 5 | How to optimize a query with LIKE? | <ul style="list-style-type: none"> ◆ LIKE '%abc%' prevents index use. ◆ LIKE 'abc%' can use index with text_pattern_ops. <input checked="" type="checkbox"/> Use full-text search or trigram index for complex patterns (PostgreSQL). |
| 6 | Handling NULL comparisons efficiently? | <ul style="list-style-type: none"> ◆ Use IS NULL, never = NULL. <input checked="" type="checkbox"/> Add DEFAULT values in schema to reduce null checks. |
| 7 | Efficient alternatives to DISTINCT? | <ul style="list-style-type: none"> ◆ DISTINCT is costly. <input checked="" type="checkbox"/> Use GROUP BY or window functions instead. ✂ Example: SELECT dept_id, MIN(emp_id) FROM employee GROUP BY dept_id; |

| No | Question | Explanation / Optimization Tip |
|----|--------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
| 8 | UNION vs UNION ALL in optimization? | ◆ UNION removes duplicates → slower . ◆ UNION ALL keeps all → faster . ✓ Use UNION ALL unless deduplication is required. |
| 9 | When to use COALESCE() in real-time queries? | ◆ Replaces NULL with default. ✓ Useful in dashboards, financial reports. ⚙ Example: SELECT emp_name, COALESCE(salary, 0) FROM employee; |
| 10 | SIMILAR TO vs LIKE? | ◆ LIKE = simple wildcards. ◆ SIMILAR TO = regex-like pattern matching. ✓ Use SIMILAR TO for complex patterns. |
| 11 | How does CASE improve report readability? | ◆ Simplifies multiple queries into one. ✓ Example: Grading logic or status columns. |
| 12 | How to optimize BETWEEN with indexes? | ◆ BETWEEN uses indexes for range queries. ✓ Ensure target column (e.g., date, salary) is indexed. |
| 13 | When should NOT IN be avoided? | ◆ NOT IN fails with NULL values. ✓ Use NOT EXISTS — it's safer and faster. |
| 14 | Optimizing multiple OR conditions? | ◆ OR degrades performance. ✓ Use IN or UNION ALL instead. ⚙ Example: WHERE dept IN ('CSE','ECE','MECH') |
| 15 | Best way to handle dynamic conditions? | ◆ Use CASE or parameters. ✓ Avoid dynamic SQL where possible. |
| 16 | Role of operator precedence in optimization? | ◆ Execution order matters: AND before OR. ✓ Always use parentheses for clarity. |
| 17 | Do operators affect query plans or index usage? | ✓ Yes. ✗ Operators like <>, NOT LIKE, NOT IN may bypass indexes . ✓ Use positive logic when possible. |
| 18 | Impact of using functions on indexed columns? | ◆ Functions (e.g., UPPER(name)) disable index . ✓ Use function-based indexes if needed. |

| No | Question | Explanation / Optimization Tip |
|----|-----------------------------------------------------------|---------------------------------------------------------------------------------|
| 19 | What operator combinations slow performance? | ◆ DISTINCT + ORDER BY + JOIN = heavy. ✓ Use CTEs or temp tables to optimize. |
| 20 | How to identify slow operator usage in PostgreSQL? | ◆ Use EXPLAIN ANALYZE. ✓ Look for Seq Scan and add index if needed. |

⚙ REAL-TIME OPTIMIZATION SCENARIOS

| Scenario | Problem | Optimization |
|--------------------------------------|--------------------------------------|----------------------------------------------------|
| Filtering large table with IN | Too many subquery values | 🔄 Replace IN with EXISTS |
| Searching with LIKE '%text%' | Index not used (full scan) | 🔍 Use GIN or Trigram index (pg_trgm in PostgreSQL) |
| Removing duplicates with DISTINCT | Memory-heavy, slow | 📊 Use GROUP BY on indexed column |
| Using NOT IN with NULLs | Missed records due to NULL logic | ⊖ Replace with NOT EXISTS |
| Repeated range queries (e.g., dates) | Slow performance | ☑ Create BTREE index on range column |
| Using functions on filter columns | Index is ignored | 🔗 Create function-based index |
| Large UNION queries | Sorting overhead | ⚡ Use UNION ALL or break logic with CTE |
| Frequent CASE conditions | Multiple queries increase complexity | 📊 Use CASE inside a single SELECT |
| Repeated NULL | Slower execution | 🌀 Use DEFAULT values in |

| Scenario | Problem | Optimization |
|----------------------------------|------------------|----------------------------------------------------------------------------------------------------------------------------------------------|
| checks | | schema to avoid NULL |
| Reporting on non-indexed columns | Full table scans |  Add composite indexes on filtered/reporting columns |

INTERVIEW AUDIO-STYLE EXPLANATION (how to speak)

Question: What's the main difference between EXISTS and IN?

Answer (speak like this):

"In SQL, both check matching conditions but in performance terms, EXISTS is more efficient for large datasets. Because it stops searching once it finds the first matching record. Whereas IN loads all subquery results into memory first. So, in real-time projects, we prefer EXISTS for optimization."

Question: When should we use UNION ALL instead of UNION?

Answer:

"If duplicates are not a problem, always use UNION ALL. It avoids sorting and de-duplication overhead. It's faster in reporting queries."

Question: Why should we avoid LIKE '%abc%' in production queries?

Answer:

"Because the leading % disables index usage and forces a sequential scan. In PostgreSQL, we use trigram or full-text search indexes to fix that."

Question: What's your approach to optimize operator-heavy queries?

Answer:

"First I check query plan using EXPLAIN ANALYZE, then identify sequential scans, and finally add or tune indexes on columns used in WHERE, JOIN, or BETWEEN clauses."

✂ Summary Table

| Operator | Common Issue | Optimization |
|-------------------|------------------------------------|----------------------------------------------------------------|
| IN / NOT IN | Slow performance with NULLs | ✓ Use EXISTS / NOT EXISTS for safer, faster logic |
| LIKE | Index not used (especially with %) | ✓ Avoid leading %; use Trigram (GIN) index (PostgreSQL) |
| DISTINCT | High memory usage on large sets | ✓ Use GROUP BY on indexed column for better performance |
| BETWEEN | May not use index | ✓ Ensure the column has a BTREE index |
| CASE | Complex, long logic in queries | ✓ Use inline logic or computed columns |
| UNION | Sorting causes performance issues | ✓ Prefer UNION ALL if duplicates aren't a concern |
| <> / NOT LIKE | Prevents index usage | ✓ Use positive conditions when possible |
| Function in WHERE | Skips index entirely | ✓ Use function-based index to restore performance |